```r
## This R script presents solutions to ECON 121 Problem Set 4.

# Clear environment, load R packages
rm(list=ls())
library(tidyverse)
library(fixest)

# Load the dataset
#load("/Users/tvogl/Dropbox/courses/econ121/data/nlsy_kids/nlsy_kids.Rdata")
load(url("https://github.com/tvogl/econ121/raw/main/data/nlsy_kids.Rdata"))
```

```r
 # # # # # #
# Problem 2 #
 # # # # # #
# Summary statistics appear below. 21 percent of the sample participated
# in HS. 32 percent of the sample is black, and 20 percent is Hispanic.
# Average mother's education is 12 years. 3 in 10 repeat a grade, another
# 3 in 10 go to college, and 7 in 10 graduate high school. Also worthy
# of note is the number of NA values, which is very high for ppvt_3.
# This high level of "missingness" will be important later.
summary(nlsy_kids)
```

```
##    head_start        sibdiff           mom_id         hispanic
##  Min.   :0.0000   Min.   :0.0000   Min.   :    3   Min.   :0.0000
##  1st Qu.:0.0000   1st Qu.:0.0000   1st Qu.: 3448   1st Qu.:0.0000
##  Median :0.0000   Median :0.0000   Median : 6400   Median :0.0000
##  Mean   :0.2066   Mean   :0.2321   Mean   : 6227   Mean   :0.2005
##  3rd Qu.:0.0000   3rd Qu.:0.0000   3rd Qu.: 8870   3rd Qu.:0.0000
##  Max.   :1.0000   Max.   :1.0000   Max.   :12667   Max.   :1.0000
##
##      black             male          firstborn        lninc_0to3
##  Min.   :0.0000   Min.   :0.0000   Min.   :0.0000   Min.   : 3.909
##  1st Qu.:0.0000   1st Qu.:0.0000   1st Qu.:0.0000   1st Qu.: 9.586
##  Median :0.0000   Median :1.0000   Median :0.0000   Median :10.118
##  Mean   :0.3203   Mean   :0.5097   Mean   :0.4045   Mean   :10.070
##  3rd Qu.:1.0000   3rd Qu.:1.0000   3rd Qu.:1.0000   3rd Qu.:10.584
##  Max.   :1.0000   Max.   :1.0000   Max.   :1.0000   Max.   :13.423
##                                                     NA's   :218
##      momed       dadhome_0to3       ppvt_3            lnbw
##  Min.   : 1.0   Min.   :0.000   Min.   :  0.00   Min.   :1.792
##  1st Qu.:10.0   1st Qu.:0.250   1st Qu.: 12.00   1st Qu.:4.635
##  Median :12.0   Median :1.000   Median : 19.00   Median :4.745
##  Mean   :11.7   Mean   :0.678   Mean   : 21.88   Mean   :4.718
##  3rd Qu.:13.0   3rd Qu.:1.000   3rd Qu.: 30.00   3rd Qu.:4.852
##  Max.   :20.0   Max.   :1.000   Max.   :101.00   Max.   :5.434
##  NA's   :6      NA's   :1603    NA's   :3591     NA's   :145
##  comp_score_5to6 comp_score_7to10 comp_score_11to14     repeat
##  Min.   : 0.00   Min.   : 0.00    Min.   : 0.6667   Min.   :0.0000
##  1st Qu.:29.50   1st Qu.:26.00    1st Qu.:23.5000   1st Qu.:0.0000
##  Median :44.50   Median :45.00    Median :42.6667   Median :0.0000
##  Mean   :45.42   Mean   :45.19    Mean   :43.7758   Mean   :0.3158
##  3rd Qu.:62.38   3rd Qu.:63.92    3rd Qu.:62.0000   3rd Qu.:1.0000
##  Max.   :98.50   Max.   :99.00    Max.   :99.0000   Max.   :1.0000
```

```
##  NA's   :1845    NA's   :1019     NA's   :1384      NA's   :1026
##     learndis           hsgrad          somecoll           idle
##  Min.   :0.00000   Min.   :0.0000   Min.   :0.0000   Min.   :0.0000
##  1st Qu.:0.00000   1st Qu.:0.0000   1st Qu.:0.0000   1st Qu.:0.0000
##  Median :0.00000   Median :1.0000   Median :0.0000   Median :0.0000
##  Mean   :0.04102   Mean   :0.7152   Mean   :0.3152   Mean   :0.1591
##  3rd Qu.:0.00000   3rd Qu.:1.0000   3rd Qu.:1.0000   3rd Qu.:0.0000
##  Max.   :1.00000   Max.   :1.0000   Max.   :1.0000   Max.   :1.0000
##  NA's   :121       NA's   :1077     NA's   :1077     NA's   :1078
##     fphealth
##  Min.   :0.0000
##  1st Qu.:0.0000
##  Median :0.0000
##  Mean   :0.0988
##  3rd Qu.:0.0000
##  Max.   :1.0000
##  NA's   :1077
```

```
# The question asks about the backgrounds of kids who participated in HS.
# HS participants are more likely to be black, have lower family income,
# and have less educated mothers, on average. They are also more likely
# to repeat a grade and less likely to go to college. However, these
# differences in long-term outcomes may reflect selection bias rather
# than the effects of HS. In other words, HS participants may have
# worse outcomes because they come from disadvantaged backgrounds.
nlsy_kids %>%
    group_by(head_start) %>%
    summarize(black = mean(black, na.rm = TRUE),
              lninc_0to3 = mean(lninc_0to3, na.rm = TRUE),
              momed = mean(momed, na.rm = TRUE),
              somecoll = mean(somecoll, na.rm = TRUE))
```

```
## # A tibble: 2 x 5
##   head_start black lninc_0to3 momed somecoll
##        <dbl> <dbl>      <dbl> <dbl>    <dbl>
## 1          0 0.269       10.1  11.8    0.329
## 2          1 0.518        9.78 11.5    0.269
```

clustering standard error

```
 # # # # # #
# Problem 3 #
 # # # # # #
# Run an OLS regression of the age 5-6 test score on the HS indicator,
# clustering standard errors by mom_id.
feols(comp_score_5to6 ~ head_start,
      data = nlsy_kids,
      vcov = ~mom_id)
```

```
## NOTE: 1,845 observations removed because of NA values (LHS: 1,845).

## OLS estimation, Dep. Var.: comp_score_5to6
## Observations: 2,420
## Standard-errors: Clustered (mom_id)
##             Estimate Std. Error  t value   Pr(>|t|)
## (Intercept) 46.65384   0.616964 75.61845  < 2.2e-16 ***
## head_start  -5.84207   1.209494 -4.83018 1.5113e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## RMSE: 22.2   Adj. R2: 0.010934
```

```
# For reference, compute the standard deviation of the test score.
sd(nlsy_kids$comp_score_5to6, na.rm=TRUE)
```

```
## [1] 22.37593
```

```
# Average scores are 5.8 points lower for participants than for non-participants.
# The association is highly statistically significant and represents roughly
# one-quarter of a standard deviation in test scores. If we assumed participation
# is exogenous, then we would conclude that HS reduces test scores by one-
# quarter of a standard deviation on average. However, we already know that
# participation is associated with several background characteristics that
# are likely to have independent effects on test scores, which implies that
# the residual is correlated with HS participation. As a result, participation
# is not exogenous, and we should not interpret the association as a causal
# effect. The bias is probably negative, since disadvantaged families select
# into HS, and kids from disadvantaged families may tend to have worse long-term
# outcomes.
```

Raw OLS

```
 # # # # # #
# Problem 4 #
 # # # # # #

# First create a data frame of families instead of kids. We can do so
# using group_by(), as follows:
nlsy_families <-
  nlsy_kids %>%
  drop_na(comp_score_5to6, head_start) %>%
  group_by(mom_id) %>%
  summarise(mean_test = mean(comp_score_5to6),
            mean_head_start = mean(head_start))

# Now estimate OLS using the family averages
feols(mean_test ~ mean_head_start,
      data = nlsy_families,
      vcov = 'hetero')
```

```
## OLS estimation, Dep. Var.: mean_test
## Observations: 1,426
## Standard-errors: Heteroskedasticity-robust
##                 Estimate Std. Error  t value   Pr(>|t|)
## (Intercept)     47.26384   0.622140 75.96982  < 2.2e-16 ***
## mean_head_start -7.58640   1.366079 -5.55341 3.3379e-08 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## RMSE: 20.0   Adj. R2: 0.018928
```

```
# The estimated coefficient on HS participation is now even more negative
# than the one from question 3. That is consistent with family-level
# omitted variables: kids from disadvantaged families enroll in HS,
# and they have have lower average test scores due to their disadvantage.
```

Fixed Effects

```
 # # # # # #
# Problem 5 #
 # # # # # #

# Estimate the model with mother fixed effects.
feols(comp_score_5to6 ~ head_start | mom_id,
      data = nlsy_kids)
```

```
## NOTE: 1,845 observations removed because of NA values (LHS: 1,845).

## OLS estimation, Dep. Var.: comp_score_5to6
## Observations: 2,420
## Fixed-effects: mom_id: 1,426
## Standard-errors: Clustered (mom_id)
##            Estimate Std. Error t value   Pr(>|t|)
## head_start  7.63285    2.01362  3.7906 0.00015655 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## RMSE: 10.7     Adj. R2: 0.442754
##              Within R2: 0.016246
```

```
# The fixed effect model suggests that HS participation raises test scores,
# in contrast to the negative effects suggested by OLS and the between effect
# model. The likely reason is that between-family variation in HS
# participation is correlated with family disadvantage, which biases us toward
# finding a negative association in the pooled and between effect models.
# The full-sample fixed effect model without controls indicates that HS
# raises test scores by 7.6 points, or one-third of a SD, on average.
```

Mother Fixed Effects Regression

```
 # # # # # #
# Problem 6 #
 # # # # # #

# In the fixed effect regression, we can include child-level covariates
# only. We cannot control for any family-level variables that do not
# vary between siblings. I choose male, firstborn, lninc_0to3,
# dadhome_0to3, and lnbw as covariates. I do not use ppvt_3 because
# it is available for few observations. When I include it, the sample
# shrinks and changes composition a lot. This was a judgment call, and
# you could have done it differently. as researchers we often face
# tradeoffs between having more information (by controlling for PPVT)
# and maintaining the composition of the sample (by not controlling for PPVT).
feols(comp_score_5to6 ~ head_start + male + firstborn + lninc_0to3 +
                        dadhome_0to3 + lnbw | mom_id,
      data = nlsy_kids)
```

```
## NOTE: 2,370 observations removed because of NA values (LHS: 1,845, RHS: 1,732).

## OLS estimation, Dep. Var.: comp_score_5to6
## Observations: 1,895
## Fixed-effects: mom_id: 1,251
## Standard-errors: Clustered (mom_id)
##              Estimate Std. Error   t value Pr(>|t|)
## head_start    5.64711    2.35257  2.400400 0.016523 *
## male         -2.81106    1.27581 -2.203352 0.027752 *
## firstborn     1.66089    1.17064  1.418783 0.156212
## lninc_0to3    2.27392    1.73535  1.310356 0.190316
## dadhome_0to3 -3.26060    3.27771 -0.994781 0.320035
## lnbw          6.91016    3.42362  2.018376 0.043765 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## RMSE: 9.48045     Adj. R2: 0.492933
##                 Within R2: 0.029561
```

```
# The estimate is still positive and statistically significant, but it
# is slightly smaller, in magnitude: HS participation raises test scores
# by 5.6 points on average. It is useful to check whether this is due to
# omitted variable bias or the different composition of the subsample
# with non-missing covariates. I re-estimate the model with no pre-HS
# covariates, but this time using the sub-sample with non-missing covariates.
# This was not necessary for full credit, but it is good practice.
nlsy_kids_subsample <-
  nlsy_kids %>%
  drop_na(male, firstborn, lninc_0to3, dadhome_0to3, lnbw)

feols(comp_score_5to6 ~ head_start | mom_id,
      data = nlsy_kids_subsample)
```

```
## NOTE: 638 observations removed because of NA values (LHS: 638).
```

```
## OLS estimation, Dep. Var.: comp_score_5to6
## Observations: 1,895
## Fixed-effects: mom_id: 1,251
## Standard-errors: Clustered (mom_id)
##            Estimate Std. Error t value Pr(>|t|)
## head_start    5.971     2.3642 2.52559 0.011673 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## RMSE: 9.5773     Adj. R2: 0.486544
##                Within R2: 0.009632
```

```
# The coefficient on HS is much closer to the regression with pre-HS
# covariates. This suggest that within-family OVB is *NOT* the issue, but
# rather that individuals with missing data on covariates have larger effects.
# The estimates are robust to controlling for pre-HS covariates
```

Standardize variable outcomes

```
 # # # # # #
# Problem 7 #
 # # # # # #

# Standardize outcome variables by subtracting mean and dividing by SD.
# The scale() function in R does this in one step:
nlsy_kids <-
  nlsy_kids %>%
  mutate(std_5to6 = scale(comp_score_5to6),
         std_7to10 = scale(comp_score_7to10),
         std_11to14 = scale(comp_score_11to14))
# You were not expected to know this function. You could have also used:
nlsy_kids <-
  nlsy_kids %>%
  mutate(stdb_5to6 = (comp_score_5to6 - mean(comp_score_5to6, na.rm = TRUE))/sd(comp_score_5to6, na.rm =
         stdb_5to6 = (comp_score_7to10 - mean(comp_score_7to10, na.rm = TRUE))/sd(comp_score_7to10, na.r
         stdb_5to6 = (comp_score_11to14 - mean(comp_score_11to14, na.rm = TRUE))/sd(comp_score_11to14, r

# Now we run a FE regression of each standardized score on HS participation,
# finding thatthe estimated effects shrink as children get older. HS raises
# scores by 0.34 standard deviations on average at ages 5-6, by 0.16 standard
# deviations at ages 7-10, and by 0.15 standard deviations at ages 11 to 14.
feols(std_5to6 ~ head_start | mom_id,
      data = nlsy_kids)
```

```
## NOTE: 1,845 observations removed because of NA values (LHS: 1,845).

## OLS estimation, Dep. Var.: std_5to6
## Observations: 2,420
## Fixed-effects: mom_id: 1,426
## Standard-errors: Clustered (mom_id)
##            Estimate Std. Error t value   Pr(>|t|)
## head_start 0.341119   0.089991  3.7906 0.00015655 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## RMSE: 0.478179     Adj. R2: 0.442754
##                  Within R2: 0.016246
```

```
feols(std_7to10 ~ head_start | mom_id,
      data = nlsy_kids)
```

```
## NOTE: 1,019 observations removed because of NA values (LHS: 1,019).

## OLS estimation, Dep. Var.: std_7to10
## Observations: 3,246
## Fixed-effects: mom_id: 1,546
## Standard-errors: Clustered (mom_id)
##            Estimate Std. Error t value Pr(>|t|)
## head_start 0.159245    0.06204 2.56682 0.010357 *
```

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## RMSE: 0.526513     Adj. R2: 0.470368
##                  Within R2: 0.004229
```

```
feols(std_11to14 ~ head_start | mom_id,
      data = nlsy_kids)
```

```
## NOTE: 1,384 observations removed because of NA values (LHS: 1,384).

## OLS estimation, Dep. Var.: std_11to14
## Observations: 2,881
## Fixed-effects: mom_id: 1,346
## Standard-errors: Clustered (mom_id)
##            Estimate Std. Error t value Pr(>|t|)
## head_start 0.153001    0.06088 2.51317 0.012081 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## RMSE: 0.511791     Adj. R2: 0.508071
##                  Within R2: 0.004263
```

```
# You may notice that the sample changes across regressions due to missingness.
# You could have also held the sample constant, as we did above for adding
# covariates. The effect on the test score at age 5-6 is still largest.
nlsy_kids_subsample <-
  nlsy_kids %>%
  drop_na(std_5to6, std_7to10, std_11to14)

feols(std_5to6 ~ head_start | mom_id,
      data = nlsy_kids_subsample)
```

```
## OLS estimation, Dep. Var.: std_5to6
## Observations: 1,728
## Fixed-effects: mom_id: 1,021
## Standard-errors: Clustered (mom_id)
##            Estimate Std. Error t value  Pr(>|t|)
## head_start 0.321301   0.103263 3.11147 0.0019133 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## RMSE: 0.472877     Adj. R2: 0.449221
##                  Within R2: 0.014944
```

```
feols(std_7to10 ~ head_start | mom_id,
      data = nlsy_kids_subsample)
```

```
## OLS estimation, Dep. Var.: std_7to10
## Observations: 1,728
## Fixed-effects: mom_id: 1,021
## Standard-errors: Clustered (mom_id)
##            Estimate Std. Error  t value Pr(>|t|)
## head_start 0.091516   0.095592 0.957356  0.33861
```

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## RMSE: 0.437615     Adj. R2: 0.536713
##                  Within R2: 0.001435
```

```
feols(std_11to14 ~ head_start | mom_id,
      data = nlsy_kids_subsample)
```

```
## OLS estimation, Dep. Var.: std_11to14
## Observations: 1,728
## Fixed-effects: mom_id: 1,021
## Standard-errors: Clustered (mom_id)
##            Estimate Std. Error t value Pr(>|t|)
## head_start 0.182914   0.101884 1.79531   0.0729 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## RMSE: 0.443297     Adj. R2: 0.524396
##                  Within R2: 0.005564
```

Final effects regression

```
 # # # # # #
# Problem 7 #
 # # # # # #

# We run FE regressions for longer-term outcomes. We find that HS participation
# reduces grade repetition by 5 percentage points, reduces learning disability
# diagnosis by 4 percentage points, raises high school graduation by 13 percentage
# points, raises college attendance by 7 percentage points, reduces idleness
# (not working or studying) by 7 percentage points, and reduces fair/poor health
# by 7 percentage points. All of these results but one (for grade repetition)
# are significant at the 5 percent level. The grade repetition result is significant
# at the 9 percent level.

feols(learndis ~ head_start | mom_id,
      data = nlsy_kids)
```

```
## NOTE: 121 observations removed because of NA values (LHS: 121).

## OLS estimation, Dep. Var.: learndis
## Observations: 4,144
## Fixed-effects: mom_id: 1,714
## Standard-errors: Clustered (mom_id)
##              Estimate Std. Error  t value  Pr(>|t|)
## head_start -0.037349   0.013224 -2.82444 0.0047912 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## RMSE: 0.144667     Adj. R2: 0.092616
##                  Within R2: 0.003505
```

```
feols(hsgrad ~ head_start | mom_id,
       data = nlsy_kids)
```

```
## NOTE: 1,077 observations removed because of NA values (LHS: 1,077).

## OLS estimation, Dep. Var.: hsgrad
## Observations: 3,188
## Fixed-effects: mom_id: 1,367
## Standard-errors: Clustered (mom_id)
##             Estimate Std. Error t value   Pr(>|t|)
## head_start 0.131179   0.030895 4.24594 2.3239e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## RMSE: 0.31008     Adj. R2: 0.17344
##                 Within R2: 0.009208
```

```
feols(somecoll ~ head_start | mom_id,
       data = nlsy_kids)
```

```
## NOTE: 1,077 observations removed because of NA values (LHS: 1,077).
```

```
## OLS estimation, Dep. Var.: somecoll
## Observations: 3,188
## Fixed-effects: mom_id: 1,367
## Standard-errors: Clustered (mom_id)
##            Estimate Std. Error t value Pr(>|t|)
## head_start 0.073996   0.030749 2.40648 0.016239 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## RMSE: 0.310531     Adj. R2: 0.217764
##                  Within R2: 0.00294
```

```
feols(idle ~ head_start | mom_id,
      data = nlsy_kids)
```

```
## NOTE: 1,078 observations removed because of NA values (LHS: 1,078).

## OLS estimation, Dep. Var.: idle
## Observations: 3,187
## Fixed-effects: mom_id: 1,367
## Standard-errors: Clustered (mom_id)
##             Estimate Std. Error  t value Pr(>|t|)
## head_start -0.072788   0.031397 -2.31828 0.020581 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## RMSE: 0.263083     Adj. R2: 0.093811
##                  Within R2: 0.003961
```

```
feols(fphealth ~ head_start | mom_id,
      data = nlsy_kids)
```

```
## NOTE: 1,077 observations removed because of NA values (LHS: 1,077).

## OLS estimation, Dep. Var.: fphealth
## Observations: 3,188
## Fixed-effects: mom_id: 1,367
## Standard-errors: Clustered (mom_id)
##             Estimate Std. Error  t value  Pr(>|t|)
## head_start -0.065942   0.023907 -2.75822 0.0058891 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## RMSE: 0.224664     Adj. R2: 0.007413
##                  Within R2: 0.004454
```

interaction terms race collinear effects

```
 # # # # # #
# Problem 8 #
 # # # # # #

# The easiest way to test for heterogeneous effects by race, ethnicity, and sex
# is include interactions of the HS dummy with race, ethnicity, and sex dummies.
# We also need to control for the main effect of sex, but not for the main effects
# or race and ethnicity because they are collinear with the mother fixed effects.
# I do this below for the high school graduation outcome. The results do not
# show strong evidence of heterogeneity in effects by race, ethnicity, or sex.
# The coefficients on the interaction terms are large, but none are significant
# at the 5% level.

# Here I use R's nice approach to interaction terms, but you could have also
# directly generated new variables for the interaction terms.

feols(hsgrad ~ head_start*(hispanic + black + male) | mom_id,
      data = nlsy_kids)
```

```
## NOTE: 1,077 observations removed because of NA values (LHS: 1,077).

## The variables 'hispanic' and 'black' have been removed because of collinearity (see $collin.var).

## OLS estimation, Dep. Var.: hsgrad
## Observations: 3,188
## Fixed-effects: mom_id: 1,367
## Standard-errors: Clustered (mom_id)
##                      Estimate Std. Error   t value   Pr(>|t|)
## head_start           0.021739   0.081351  0.267217 7.8934e-01
## male                -0.108291   0.022182 -4.882017 1.1737e-06 ***
## head_start:hispanic  0.071468   0.097529  0.732784 4.6382e-01
## head_start:black     0.110709   0.087847  1.260241 2.0780e-01
## head_start:male      0.062209   0.045927  1.354534 1.7579e-01
## ... 2 variables were removed because of collinearity (hispanic and black)
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## RMSE: 0.307758     Adj. R2: 0.18398
##                  Within R2: 0.023991
```

conclusion

```
 # # # # # #
# Problem 9 #
 # # # # # #

# The evidence suggests that HS participation has lasting effects on children's
# outcomes, which provides some justification for the program's existence. Whether
# the government chould expand or cut funding for this and similar programs depends
# on its cost-effectiveness compared with other potential use of funds. In general,
# it is difficult to extrapolate the effects of program expansion from our estimated
# average effects of treatment on the treated because the effects may be different
# in the new subpopulations that would gain access if the program expanded. At the
# same time, the lack of significant treatment effect heterogeneity in Problem 9
# suggests that perhaps we can extrapolate. Many answers could receive full credit
# for this question.
```